

What has been built

A progress report against Build Specification v2.0 of 20 August 2026 — its 17 clauses, 116 numbered requirements and 12 component tests. **Every clause is built and deployed.** Nothing in the system is currently known to be broken. Nothing has been accepted either, and this document is as specific about the second as the first.

17 clauses built and audited	116 numbered requirements	95 automated tests passing	0/12 acceptance tests met yet
------------------------------------	---------------------------------	----------------------------------	-------------------------------------

What you can use today

The Telegram bot takes voice notes, interviews you, and runs your weekly pass — type `/` and it lists what it can do. The desk is live at **the-desk-saip00519-gmailcoms-projects.vercel.app**, holding the idea bank, the calendar, the reference library and the change proposals, with $\aleph K$ search across everything. The database runs itself on nine schedules.

What is honestly not ready

No draft has yet cleared the quality gate, because no paid model has been configured. LinkedIn is not connected. No image has been rebuilt. Seven of fourteen reference library sections are waiting on you, and the voice guide is the one that matters.

Every figure in this document was read from the live system on the date above, with test fixtures excluded, and can be reproduced by running `node eval/acceptance.mjs`.

A system that turns things you already say — voice notes, calls, working sessions — into LinkedIn posts that **could only have been written by you**. Three rules were never negotiable. Each is now enforced by the database rather than by our good intentions, which means they hold even if the code above them is wrong.

The lived experience test

clause 1

ENFORCED

Could anyone else have written this? If yes, it fails. Held in a locked section of your reference library. The learning loop can propose changes to anything else and is **structurally unable to touch this** — a database rule rejects the attempt.

No fabrication

9.4

ENFORCED

Every claim, quote, number and name must trace to something you actually said. The writer emits each claim with the exact source text beside it, and that is checked **mechanically, before any model gives an opinion**. Nothing invented survives.

Nothing publishes without you

11.2

ENFORCED

Four database constraints make an unauthorised published post **impossible to store**, not merely unlikely. Tested adversarially: the database refuses a published post you never marked ready, one with no date, and one published before you approved it.

The whole pipeline exists end to end. **Proven** means it has been driven against the live system and watched working, not read in the code. **Unproven** means it is built and waits on something outside our control.

Five ways in

clause 4

PROVEN

Voice notes, prompted sessions, call transcripts, Claude Code sessions and typed thoughts, all landing in one idea bank. Send a voice note from your phone and it is transcribed, filed and answered with a question — in seconds, between other things. **9 inputs captured so far.**

The interview

clause 5

PROVEN

Asks until a moment has a scene, a detail and a realisation — then stops. It budgets its questions, rewords prompts that never land, and **parks a moment rather than padding it**. It never writes; that is a different job with a different prompt.

The idea bank

clause 6

PROVEN

28 tables, 73 access rules, and **no delete permission anywhere** — nothing can be destroyed, only killed and labelled. Moments re-open when you add to them. Editable by you without a developer, and exportable in full at any moment.

Choosing what gets written

clause 7

PROVEN

Five-factor scoring, pillar balance, and a retelling check against everything already published. When it cannot fill the queue without dropping the bar it **runs short and says why**, because fewer posts is the right answer when the material is not there. **46 selection runs.**

Drafting and the gate

clause 9

BUILT, UNPROVEN

Eight independent checks, each able to reject on its own — because a model asked to weigh eight things at once averages across them and lets marginal work through. **130 checks have run and 48 rejected.** The gate demonstrably rejects. What it has never done is pass a draft, and section 4 explains exactly why.

Visuals

clause 10

BUILT, UNPROVEN

Send a diagram and it rebuilds the *idea* in your brand — drawn from scratch as a vector image, so it is **structurally incapable of resembling the original** rather than merely instructed not to. Your image is stored and queued; no rebuild has run yet.

The calendar

clause 11

PROVEN

Drafts arrive as drafts. You edit in place and mark ready, and only that authorises publishing. Your edits are read back into the bank automatically, because the calendar **is** the bank rather than a copy of it.

The learning loop

clause 12

PROVEN

The difference between what was written and what you approved is measured **the moment you approve** — it needs no LinkedIn and no effort from you. A one-line verdict is asked on the taps you already make. Changes to the library are proposed with evidence, approved only by you, and reversible.

Operations

clause 13

PROVEN

Silence alerts, queue warnings, errors surfaced once rather than every morning, and a monthly report on the system itself. It also tells you if it starts using a service you were never told about. **Nine schedules run without anyone starting them.**

Four things were built. Two you use directly, one runs quietly on your machine, and one is the engine underneath. **The split between the first two is deliberate and worth understanding**, because it decides where everything happens.

The Telegram bot — every decision

LIVE

Where the system is *used*. Capture, the interview, the weekly pass, approving posts, clearing names, sending images, adding a rule. **Type / and it lists everything it can do** — nine commands, each described by what it does for you. `/help` and `/status` answer with buttons that run the thing rather than naming it, and the buttons change with what is actually waiting: an empty queue offers you a session, a full one offers the review.

Every decision the system needs from you can be made with a tap. Nothing here requires a laptop, because 12.7 says a habit that needs a second place to go will not survive.

The desk — the four jobs a chat box cannot do

LIVE

Deliberately *not* a second copy of the bot. Clause 6.1 requires a surface that is not the chat interface, and this is the small number of jobs that genuinely need a keyboard and a screen: **editing a post, writing the reference library, finding a moment, and rolling a change back or taking your data out**.

Six screens and an export. **⌘K searches every moment** — your words, the scene, the detail, what they actually said — ranked, across every status, because the moment you go looking for later is usually one that has already moved on. Built as proper full-text search rather than pattern matching, so it finds "costs" when you type "cost" and does not return "costume".

The Claude Code agent and its installer

BUILT, NOT INSTALLED

Reads your working sessions on your own machine and sends a **digest of what you said and decided** — never code, never client material. That split is not politeness: 625 MB of logs would be unaffordable to send, and 15.4 forbids putting client work into a third-party context. It filters hard locally and stays silent about routine work.

One setup command puts it on a schedule and it never needs attention again. `--doctor` checks every reason it could fail, including whether the credentials actually work against the live server — **because a wrong key produces silence, and a quiet week is a legitimate pass here**, so silence and breakage would otherwise look identical. It has not been installed yet, so this input has produced nothing.

The engine

LIVE

10 background services and 27 database migrations, running on nine schedules — the queue every minute, selection every four hours, publishing every five minutes, the daily check at 8am, the learning run on Monday mornings. **285 model calls so far, at a total cost of \$0.09**.

ALSO DELIVERED

Six written documents, because clause 14 requires them and because you should not need us to run this: the spec read clause by clause, the seven build decisions with the reasoning behind each, the running cost, the operating runbook, the account handover sequence, and the running order for the recorded walkthrough. Plus **95 automated tests** and four evaluation harnesses that measure the system against its own acceptance criteria.

IF SOMETHING HAD TO BE CUT

It would be the Claude Code reader. It is one of five ways in, it is the only one needing software on your laptop, and by design it produces nothing in a quiet week. Its value is that it catches the material you generate without noticing — the session where an assumption turned out to be wrong. That is real "only you were in that room" material, but it is a bonus input rather than a load-bearing one, and nothing else in the build depends on it.

Twelve component tests decide whether this is finished. **None has been passed and none has been failed** — every one is waiting on something specific, and this is that list.

#	TEST	WAITING ON	WHERE IT STANDS
01	Raw capture	You	1 of 20 voice notes sent. The mechanism is proven; the volume is not there.
02	Prompted session	You	3 of 5 sessions run.
03	Call transcripts	You, then 2 weeks	Needs a call recorder chosen, then two weeks of calls.
04	Claude Code	One install	The agent has never run. One command starts a two-week clock.
05	Interview	You	1 of 2 answered moments reached depth. Five more were asked and never answered — which measures nothing about the interview.
06	Fabrication	A paid model	Needs 10 drafts through the gate for you to audit. There are none.
07	Names	A paid model	Same. Four client names are on file and correctly held back.
08	The gate	A paid model	Run for real: 6 of 80 checks completed before the free model hit its limit. Not a failure — it never got to read them.
09	Selection	4 weeks	46 runs so far. Needs four weeks of elapsed running.
10	Visuals	A paid model	Needs a model that can see pictures. Your image is stored and queued.
11	Calendar	LinkedIn	Zero exceptions, but the four-week window has not begun. The guarantee itself is proven.
12	Learning loop	LinkedIn, 1 month	The mechanism runs. Nothing has published for it to learn from.
17a	The finish line	All of the above	Five of the last six drafts needing only light editing. No draft has been approved yet, so there is nothing to measure — reported as unmeasurable rather than as zero.

THE ONE THING THAT UNBLOCKS THE MOST

A paid model key. The build has been running on free tiers, which allow 8,000 words a minute where a single quality check needs roughly 16,000. That one fact blocks four component tests and the finish line itself.

It is worth being precise about what this is *not*: the gate is not broken. Reading its rejections shows it correctly catching genuine invention — phrases the free model added that you never said. **The safety mechanism works; the writer beneath it is the cheap one.** This is already inside the \$53 a month confirmed in the cost document.

Seven of the fourteen reference library sections are waiting on you, and they are the ones that cannot be guessed without putting words in your mouth.

The voice guide

8.1, 8.2

BLOCKING

A recording of you talking at length, in your own words — **never your old posts**, which you have said have drifted. The entire quality argument rests on this, and it has no date against it. The system has already written you a first draft to react to, which is a far easier thing to answer than a blank page. It needs about 400 words of you speaking first.

Pillars, audience, gate rules, visual brand, reference posts

clause 8

WAITING

Each section explains what belongs in it and why we did not draft it for you. Editable from The desk in a browser, or from Telegram with a single message — a rule can be added in under a minute, because that is how it will actually get maintained.

Connect LinkedIn

clause 11, 12.1

WAITING

Posting is self-serve and needs no approval. Post analytics sits behind a review that takes time, so it is worth starting early. This begins the four-week clocks on two component tests.

Answer what is already asked, and keep sending

clauses 4, 5

WAITING

Five moments are sitting asked and unanswered. Beyond those: nineteen more voice notes, two more sessions, and a decision on which call recorder you use.

THE HONEST CONSTRAINT

The critical path is calendar time, not code. Five of the twelve tests need two to four weeks of live running, and those clocks have not started. No amount of building shortens them. The most valuable thing anyone can do this week is start one.

It becomes yours

All code, prompts, configuration and documentation transfer to you on payment, with no ongoing licence from us. Every third-party component is permissively licensed and listed by name, so *"no licence required"* is checkable rather than asserted.

It runs without us

Nothing in the design depends on Thought Pilot. The accounts currently do, which is **the one clause not yet met** — everything runs today in our Supabase organisation, on our keys. The move into your accounts has been rehearsed rather than described: all 26 migrations were applied to an empty database to prove the sequence works.

What it costs

\$45–70 a month, typically \$53 — against the roughly \$50 you expected. Confirmed in writing before the build, itemised line by line, with three levers if you want it firmly under \$50.

Today it is running at close to zero, on free tiers, which is exactly why the gate has not been cleared. Two decisions cost money and we would not economise on either: the eight-check gate at about \$6, and database backups at \$25.

You will be told before a bill

The system now watches which services it actually uses and tells you in the daily message the first time it touches one that is not on your list.

WHERE IT RUNS TODAY

The desk is at `the-desk-saip00519-gmailcoms-projects.vercel.app`, hosted free, running in Sydney beside the database so a page load does not cross an ocean. Sign-in is restricted to one address and the door is locked by default rather than open by default.

The database is in Sydney on Supabase; the Telegram bot needs no hosting of its own. All of it is currently in **Thought Pilot accounts**, which is the one clause not yet met and the first thing to change at handover.

TWO THINGS WORTH RAISING DIRECTLY

Two services were added mid-build without telling you. Groq and Hugging Face have been running the writer, the gate and the interview on free tiers while the paid key was not configured, and the tool list you hold named neither. That was a defensible decision made silently, which is not good enough. Both are now on the list, and the system was changed so it cannot happen quietly again.

Sign-in was open when it should have been closed. The single-address restriction had a fault that let anyone signed in reach the app if the setting was missing — and the setting had never been filled in. Found while deploying, fixed before the app went public, and no real client material has ever been in the system. Development keys will be rotated before any is.